

Improved Partial Effects Plots for GAMs Looking at the Plankton Community

Curtis C. Bohlen, Casco Bay Estuary Partnership

7/21/2022

Contents

Introduction	2
General Instructions to Authors About Graphics	2
Load Libraries	2
Set Graphics Theme	3
Input Data	3
Folder References	3
Load Data	3
Complete Cases	5
Reduced Data	6
Functions for Graphics Preparation	6
Find Evenly Spaced Points	6
Conduct The Analysis	6
Construct a Plot	7
Total Zooplankton Density	8
Reduced Complexity Model	8
Generate Separate Marginal Means	9
Combine Marginal Means Data	9
Duplicate the Raw Data	9
Save the Plot	11
Shannon Diversity	11
Model on Reduced Data	11

Single Species Models	14
Model Choice	14
Automating Analysis of Separate Species	14
Acartia	15
Balanus	17

Introduction

This notebook reprises selected analyses using GAMs, and then develops nicer partial effects plots than the `mgcv` defaults.

In particular, I'm interested in generating marginal plots as follows (the page numbers refer to where the “draft” marginal plots appear in the PDF of the original analysis).

1. Zoop density vs. Turbidity (pg. 20)
2. Zoop diversity vs. Chl & Zoop diversity vs. Turbidity (combined 2-part figure), pg. 25
3. Eurytemora density vs. Turbidity, pg. 34
4. Barnacle vs. Chl & Barnacle vs. Temp (combined 2-part figure), pg. 31
5. Acartia vs. Temp & Acartia vs. Salinity (combined 2-part figure), pg. 29

Rachel Lasley-Rasher has requested that I assemble plots that lack repeated y axes. It is probably possible to construct them directly, assembling a graphic out of two GROBs containing graphs and a third containing the axis title and axis labels, but that would not be easy.

A possible alternative is to assemble artificial data frames specifically to facilitate use of `ggplot2`'s faceting capabilities. That is the approach I take here.

The core idea is that you can assemble a couple of synthetic data frames (one for marginal means, one for raw data) to allow you to facet the resulting graphics, when you have multiple graphics. This function should not have to be changed, since you can facet the output of this function.

General Instructions to Authors About Graphics

The instructions to authors suggests figure widths should line up with columns, and proposes figure widths should be:

39 mm ~ 1.54 inches 84 mm ~ 3.30 inches 129 mm ~ 5.04 inches 174 mm ~ 6.85 inches

With height not to exceed 235 mm (9.25 inches).

RMarkdown / `knitr` likes figure dimensions in inches. 174 mm is about 6.85 inches

Load Libraries

```
library(tidyverse)
#> -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
#> v dplyr      1.1.1      v readr      2.1.4
#> v forcats    1.0.0      v stringr    1.5.0
#> v ggplot2     3.4.1      v tibble     3.2.1
```

```

#> v lubridate 1.9.2      v tidyr      1.3.0
#> v purrr      1.0.1
#> -- Conflicts ----- tidyverse_conflicts() --
#> x dplyr::filter() masks stats::filter()
#> x dplyr::lag()     masks stats::lag()
#> i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
library(readxl)
library(mgcv)      # for GAM models
#> Loading required package: nlme
#>
#> Attaching package: 'nlme'
#>
#> The following object is masked from 'package:dplyr':
#>
#> collapse
#>
#> This is mgcv 1.8-42. For overview type 'help("mgcv-package")'.
library(emmeans)
library(ggeffects)
library(gridExtra) # or could use related functions in `cowplot`
#>
#> Attaching package: 'gridExtra'
#>
#> The following object is masked from 'package:dplyr':
#>
#> combine

```

Set Graphics Theme

This sets `ggplot()` graphics for no background, no grid lines, etc. in a clean format suitable for (some) publications.

```
theme_set(theme_classic())
```

Input Data

Folder References

```

data_folder <- "Original_Data"

dir.create(file.path(getwd(), 'figures'), showWarnings = FALSE)

```

Load Data

```

filename.in <- "penob.station.data EA 3.12.20.xlsx"
file_path <- file.path(data_folder, filename.in)
station_data <- read_excel(file_path,

```

```

        sheet="Final", col_types = c("skip", "date",
                                     "numeric", "text", "numeric",
                                     "text", "skip", "skip",
                                     "skip",
                                     rep("numeric", 10),
                                     "text",
                                     rep("numeric", 47),
                                     "text",
                                     rep("numeric", 12))) %>%

rename_with(~ gsub(" ", "_", .x)) %>%
rename_with(~ gsub("\\.", "_", .x)) %>%
rename_with(~ gsub("\\?", "", .x)) %>%
rename_with(~ gsub("%", "pct", .x)) %>%
rename_with(~ gsub("_Abundance", "", .x)) %>%
filter(! is.na(date))
#> New names:
#> * `` -> `...61`

```

```

names(station_data)[10:12]
#> [1] "discharge_week_cftpersec" "discharg_day"
#> [3] "discharge_week_max"
names(station_data)[10:12] <- c('disch_wk', 'disch_day', 'disch_max')

```

Station names are arbitrary, and Erin previously expressed interest in renaming them from Stations 2, 4, 5 and 8 to Stations 1,2,3,and 4.

The `factor()` function by default sorts levels before assigning numeric codes, so a convenient way to replace the existing station codes with sequential numbers is to create a factor and extract the numeric indicator values with `as.numeric()`.

```

station_data <- station_data %>%
  mutate(station = factor(as.numeric(factor(station))))
head(station_data)
#> # A tibble: 6 x 76
#>   date          year month month_num season riv_km station station_num
#>   <dtm>         <dbl> <chr>      <dbl> <chr>   <dbl> <fct>      <dbl>
#> 1 2013-05-28 00:00:00 2013 May          5 Spring  22.6  1          1
#> 2 2013-05-28 00:00:00 2013 May          5 Spring  13.9  2          2
#> 3 2013-05-28 00:00:00 2013 May          5 Spring   8.12 3          3
#> 4 2013-05-28 00:00:00 2013 May          5 Spring   2.78 4          4
#> 5 2013-07-25 00:00:00 2013 July          7 Summer  22.6  1          1
#> 6 2013-07-25 00:00:00 2013 July          7 Summer  13.9  2          2
#> # i 68 more variables: depth <dbl>, disch_wk <dbl>, disch_day <dbl>,
#> #   disch_max <dbl>, tide_height <dbl>, Full_Moon <dbl>, Abs_Moon <dbl>,
#> #   Spring_or_Neap <chr>, ave_temp_c <dbl>, ave_sal_psu <dbl>,
#> #   ave_turb_ntu <dbl>, ave_do_mgperl <dbl>, ave_DO_Saturation <dbl>,
#> #   ave_chl_microgperl <dbl>, sur_temp <dbl>, sur_sal <dbl>, sur_turb <dbl>,
#> #   sur_do <dbl>, sur_chl <dbl>, bot_temp <dbl>, bot_sal <dbl>, bot_turb <dbl>,
#> #   bot_do <dbl>, bot_chl <dbl>, max_temp <dbl>, max_sal <dbl>, ...

```

Subsetting to Desired Data Columns

I base selection of predictor variables here on the ones used in the manuscript.

```

base_data <- station_data %>%
  rename(Date = date,
         Station = station,
         Year = year) %>%
  select(-c(month, month_num)) %>%
  mutate(Month = factor(as.numeric(format(Date, format = '%m')),
                       levels = 1:12,
                       labels = month.abb),

         DOY = as.numeric(format(Date, format = '%j')),
         season = factor(season, levels = c('Spring', 'Summer', 'Fall')),
         is_sp_up = season == 'Spring' & Station == 1,
         Yearf = factor(Year)) %>%
  rename(Season = season,
         Density = combined_density,
         Temp = ave_temp_c,
         Sal = ave_sal_psu,
         Turb = sur_turb,
         AvgTurb = ave_turb_ntu,
         DOSat = ave_DO_Saturation,
         Chl = ave_chl_microgperl,
         Fish = `___61`,
         RH = Herring
  ) %>%
  select(Date, Station, Year, Yearf, Month, Season, is_sp_up, DOY, riv_km,
         disch_wk, disch_day, disch_max,
         Temp, Sal, Turb, AvgTurb, DOSat, Chl,
         Fish, RH,
         Density, H, SEI,
         Acartia, Balanus, Eurytemora, Polychaete, Pseudocal, Temora) %>%
  arrange(Date, Station)
head(base_data)
#> # A tibble: 6 x 29
#>   Date           Station Year Yearf Month Season is_sp_up DOY riv_km
#>   <dtm>         <fct>   <dbl> <fct> <fct> <fct> <lgl>   <dbl> <dbl>
#> 1 2013-05-28 00:00:00 1      2013 2013 May   Spring TRUE    148  22.6
#> 2 2013-05-28 00:00:00 2      2013 2013 May   Spring FALSE   148  13.9
#> 3 2013-05-28 00:00:00 3      2013 2013 May   Spring FALSE   148   8.12
#> 4 2013-05-28 00:00:00 4      2013 2013 May   Spring FALSE   148   2.78
#> 5 2013-07-25 00:00:00 1      2013 2013 Jul    Summer FALSE   206  22.6
#> 6 2013-07-25 00:00:00 2      2013 2013 Jul    Summer FALSE   206  13.9
#> # i 20 more variables: disch_wk <dbl>, disch_day <dbl>, disch_max <dbl>,
#> #   Temp <dbl>, Sal <dbl>, Turb <dbl>, AvgTurb <dbl>, DOSat <dbl>, Chl <dbl>,
#> #   Fish <dbl>, RH <dbl>, Density <dbl>, H <dbl>, SEI <dbl>, Acartia <dbl>,
#> #   Balanus <dbl>, Eurytemora <dbl>, Polychaete <dbl>, Pseudocal <dbl>,
#> #   Temora <dbl>

```

```
rm(station_data)
```

Complete Cases

This drops only two samples, one for missing Zooplankton data, one for missing fish data. We need this reduced data set to run The `step()` function. It makes little sense to try stepwise model selection if each

time you add or remove a variable, the sample you are studying changes. Since fish is never an important predictor, we will want need to refit models after stepwise elimination to use the most complete possible data set.

```
complete_data <- base_data %>%
  select(Season, Station, Yearf,
         is_sp_up, Temp, Sal, Turb, Chl, Fish, RH,
         Density, H,
         Acartia, Balanus, Eurytemora, Polychaete, Pseudocal, Temora) %>%
  filter(complete.cases(.))
```

Reduced Data

The low salinity spring samples are doing something rather different, and they complicate model fitting. Models are far better behaved if we exclude a few extreme samples. These are low salinity low zooplankton samples. We have two complementary ways to specify which samples to omit, without just omitting “outliers”. The first is to restrict modeling to “marine” samples over a certain salinity range, and the other is to omit spring upstream samples, which include most of the problematic samples. We eventually decided to go with the first.

```
drop_low <- complete_data %>%
  filter(Sal > 10)      # Pulls three samples, including one fall upstream sample
                        # a fourth low salinity sample lacks zooplankton data
#drop_sp_up <- complete_data %>%
# filter(! is_sp_up) # drops four samples
```

Functions for Graphics Preparation

I developed versions of these functions in the “GAM Analysis Partials.Rmd” notebook. See that notebook and the “Testing indirection.Rmd” notebook for more of the logic involved.

Find Evenly Spaced Points

This finds evenly spaced points along the range of a specified variable.

```
find_stops <- function(.dat, .predictor, .nstops = 25) {
  .predictor <- ensym(.predictor)
  r <- range(.dat[.[.predictor]])
  stops = seq(r[1], r[2], length.out = .nstops)
  return(stops)
}
```

Conduct The Analysis

This calculates marginal means along one predictor variable in a model. Much of the code complexity handles special cases where either the x or y variables are transformed, which changes the way parts of the output are named.

```

marginal_analysis <- function(.dat, .predictor, .model,
                             .nstops = 25, .logx = TRUE, .transy = TRUE) {
  .predictor <- ensym(.predictor)

  the_name <- as.character(.predictor)
  the_log_name <- paste0("log(", the_name, ")")

  # The following finds stops linear in the original predictor scale.
  # That is appropriate for the planned graphics, where both axes are
  # untransformed.
  stops <- find_stops(.dat, !!.predictor, .nstops)
  # browser()
  if (.logx) {
    stopslist <- list(log(stops))
    names(stopslist) <- the_log_name

    emms <- emmeans(.model, the_log_name,
                    at = stopslist,
                    type = 'response')
    emms <- as_tibble(emms)
    #browser()
    emms <- emms %>%
      mutate(!!the_name := exp(emms[[the_log_name]]))
  }
  else {
    #browser()
    stopslist <- list(stops)
    names(stopslist) <- the_name

    emms <- emmeans(.model, the_name,
                    at = stopslist,
                    type = 'response')
    emms <- as_tibble(emms)
  }
  #The default name of the output of emmeans() differs if the response
  #variable is transformed or untransformed. This makes names consistent.
  if (!.transy) {
    emms <- emms %>%
      rename(response = emmean)
  }
  return(emms)
}

```

Construct a Plot

This function does not handle axis labels (to simplify the code).

```

marginal_plot <- function(.emms, .data, .predictor, .response, .panel = NULL) {
  browser()
  .predictor <- ensym(.predictor)
  .response <- ensym(.response)

```

```

the_name      <- as.character(.predictor)
the_log_name  <- paste0("log(", the_name, ")")

xlocs <- .emms[[the_name]]

#"response" is the default name of the output column in `emmeans()`
# lower.CL and # upper.CL are also default column names.

ggplot(.emms, aes(xlocs, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = .data, mapping = aes(x = !!.predictor, y = !!.response),
            size = 1, color = "gray40") +
  geom_rug(data = .data, mapping = aes(x = !!.predictor, y = NULL)) +
  theme(axis.title = element_text(size = 9),
        axis.text = element_text(size = 8))
}

```

Total Zooplankton Density

I fit the simplified model without Station. The full model has the same concavity problems as before, and here the model fails to converge. While I could alter the convergence criteria to search for a solution, we know the model that includes Station will have concavity problems, so there is little point.

Reduced Complexity Model

```

density_gam_reduced <- gam(log(Density) ~
  #s(Temp, bs="ts", k = 5) +
  #s(Sal, bs="ts", k = 5) +
  s(log(Turb), bs="ts", k = 5) +
  s(log(Chl), bs="ts", k = 5) +
  #s(log1p(Fish), bs="ts", k = 5) +
  s(Yearf, bs = 're'),
  data = drop_low, family = 'gaussian')
summary(density_gam_reduced)
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> log(Density) ~ s(log(Turb), bs = "ts", k = 5) + s(log(Chl), bs = "ts",
#>      k = 5) + s(Yearf, bs = "re")
#>
#> Parametric coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   8.1283      0.2307   35.23  <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:

```



```
#>               edf Ref.df      F p-value
#> s(log(Turb)) 1.4120      4  6.18 0.000253 ***
#> s(log(Chl))  0.6072      4  0.83 0.122462
#> s(Yearf)     3.6720      4 10.52 1.63e-06 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) = 0.561   Deviance explained = 60.8%
#> GCV = 0.26018   Scale est. = 0.22853   n = 55
```

Generate Separate Marginal Means

```
turb <- marginal_analysis(drop_low, Turb, density_gam_reduced,
                          .nstops = 25, .logx = TRUE)
```

```
chl <- marginal_analysis(drop_low, Chl, density_gam_reduced,
                        .nstops = 25, .logx = TRUE)
```

Combine Marginal Means Data

Now, we need to assemble these into a single data frame, with appropriate labels to clarify which variable they represent. Also the predictor variables have different names, so we need to harmonize them.

```
names(turb) <- c("log(Pred)", names(turb)[2:6], "Pred")
names(chl) <- c("log(Pred)", names(chl)[2:6], "Pred")
```

```
emms <- bind_rows(Turbidity = turb, Chlorophyll = chl, .id = "source")
```

Duplicate the Raw Data

With appropriate names.

```
dat <- drop_low %>%
  select(Turb, Chl, Density) %>%
  pivot_longer(c(Turb, Chl), names_to = "source", values_to = "Pred") %>%
  mutate(source = fct_recode(source, Turbidity = "Turb", Chlorophyll = "Chl"))
```

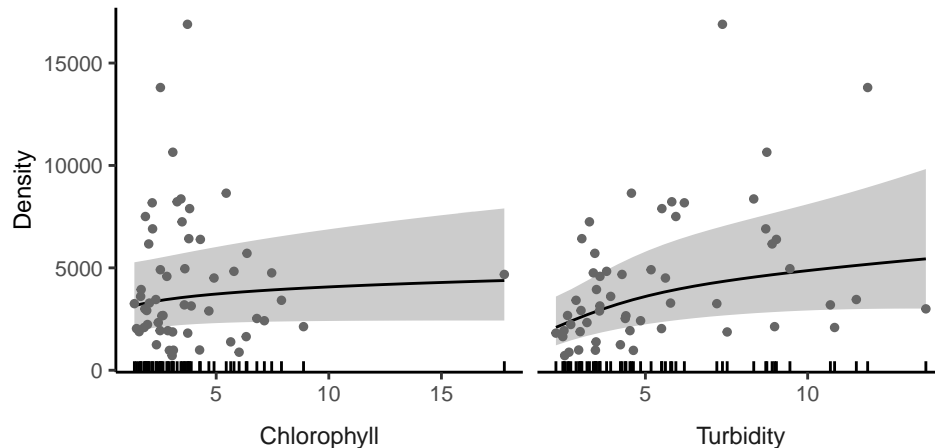
```
ggplot(emms, aes(Pred, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = dat, mapping = aes(x = Pred, y = Density),
            size = 1, color = "grey40") +
  geom_rug(data = dat, mapping = aes(x = Pred, y = NULL)) +
  facet_wrap(~source, scales = "free_x", strip.position = "bottom") +

  theme(axis.title = element_text(size = 9),
        axis.text = element_text(size = 8),
        axis.title.x = element_blank(),
```

```

strip.background = element_blank(),
strip.placement = "outside",
strip.text = element_text(size = 9)) +
ylab("Density")

```



Pretty close, but now we need to get the labels to show as fancy labels built up out of expressions.

```

fancy_turb <- expression("Turbidity" ~ "(NTU)")
fancy_chl  <- expression("Chlorophyll (" * mu * g * L ^-1 ~ ")")
fancy_dens <- expression("Zooplankton Density (" * m ^-3 ~ ")")
dat <- drop_low %>%
  select(Turb, Chl, Density) %>%
  pivot_longer(c(Turb, Chl), names_to = "source", values_to = "Pred") %>%
  mutate(source = factor(source,
    levels = c("Turb", "Chl"),
    labels = c(fancy_turb, fancy_chl)))

emms <- bind_rows(Turbidity = turb, Chlorophyll = chl, .id = "source") %>%
  mutate(source = factor(source,
    levels = c("Turbidity", "Chlorophyll"),
    labels = c(fancy_turb, fancy_chl)))

```

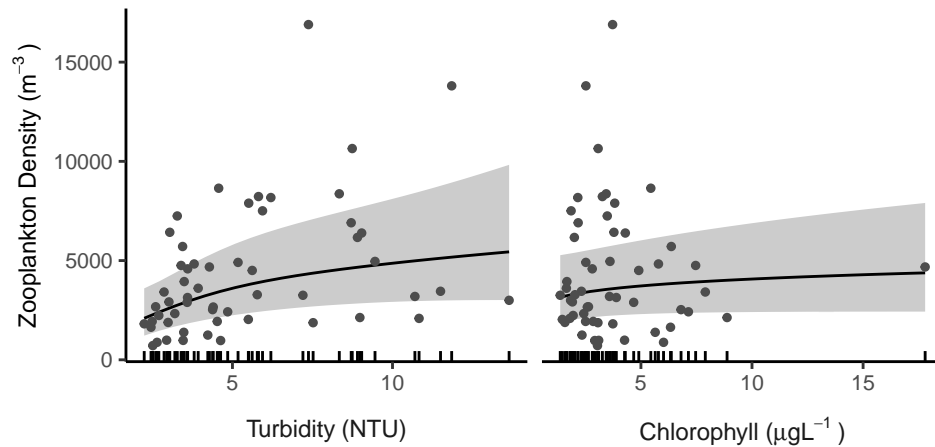
```

ggplot(emms, aes(Pred, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = dat, mapping = aes(x = Pred, y = Density),
    size = 1, color = "grey30") +
  geom_rug(data = dat, mapping = aes(x = Pred, y = NULL)) +
  facet_wrap(~source, scales = "free_x", strip.position = "bottom",
    labeller=label_parsed) +

  theme(axis.title = element_text(size = 9),
    axis.text = element_text(size = 8),
    axis.title.x = element_blank(),
    strip.background = element_blank(),
    strip.placement = "outside",

```

```
strip.text = element_text(size = 9)) +
ylab(fancy_dens)
```



Save the Plot

```
ggsave(file='figures/density_2.png',
width = 5.04, height = 2.2)
ggsave('figures/density_2.pdf', device = cairo_pdf,
width = 5.04, height = 2.2)
```

Shannon Diversity

Model on Reduced Data

```
shannon_gam_no_low <- gam(H ~
  s(Temp, bs="ts", k = 5) +
  s(Sal, bs="ts", k = 5) +
  s(log(Turb), bs="ts", k = 5) +
  s(log(Chl), bs="ts", k = 5) +
  s(log1p(Fish),bs="ts", k = 5) +
  s(Yearf, bs = 're'),
  data = drop_low, family = 'gaussian')
summary(shannon_gam_no_low)
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> H ~ s(Temp, bs = "ts", k = 5) + s(Sal, bs = "ts", k = 5) + s(log(Turb),
#>   bs = "ts", k = 5) + s(log(Chl), bs = "ts", k = 5) + s(log1p(Fish),
#>   bs = "ts", k = 5) + s(Yearf, bs = "re")
#>
```

```

#> Parametric coefficients:
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    1.3310      0.1142   11.66  3.1e-15 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:
#>               edf Ref.df       F  p-value
#> s(Temp)         1.615e+00     4  4.222 0.002901 **
#> s(Sal)           2.259e-08     4  0.000 0.257386
#> s(log(Turb))     1.369e-08     4  0.000 0.608480
#> s(log(Chl))      3.767e+00     4 11.002 0.000252 ***
#> s(log1p(Fish))   3.675e-01     4  0.167 0.197576
#> s(Yearf)         2.929e+00     4  2.802 0.008131 **
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) =  0.417   Deviance explained = 51.1%
#> GCV =      0.2   Scale est. = 0.1648      n = 55

```

Generate Separate Marginal Means

```

temp <- marginal_analysis(drop_low, Temp, shannon_gam_no_low,
                          .nstops = 25, .logx = FALSE)

```

```

chl <- marginal_analysis(drop_low, Chl, shannon_gam_no_low,
                        .nstops = 25, .logx = TRUE, .transy = FALSE)

```

```

names(temp) <- c("Pred", names(temp)[2:6])
names(chl) <- c("log(Pred)", names(chl)[2:6], "Pred")
chl <- chl[,c(7, 2:6)]

```

Name match

Assemble data

```

fancy_temp <- expression("Temperature (" * degree * "C)")
fancy_chl  <- expression("Chlorophyll (" * mu * g * L ^-1 ~")")
dat <- drop_low %>%
  select(Temp, Chl, H) %>%
  pivot_longer(c(Temp, Chl), names_to = "source", values_to = "Pred") %>%
  mutate(source = factor(source,
                        levels = c("Temp", "Chl"),
                        labels = c(fancy_temp, fancy_chl)))

emms <- bind_rows(Temperature = temp, Chlorophyll = chl, .id = "source") %>%

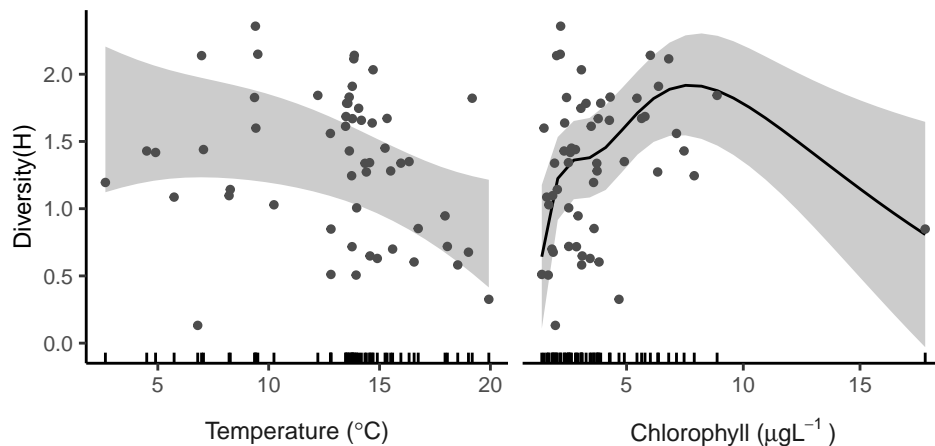
```

```
mutate(source = factor(source,
                        levels = c("Temperature", "Chlorophyll"),
                        labels = c(fancy_temp, fancy_chl)))
```

Generate Plot

```
ggplot(emms, aes(Pred, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = dat, mapping = aes(x = Pred, y = H),
            size = 1, color = "grey30") +
  geom_rug(data = dat, mapping = aes(x = Pred, y = NULL)) +
  facet_wrap(~source, scales = "free_x", strip.position = "bottom",
            labeller=label_parsed) +

  theme(axis.title = element_text(size = 9),
        axis.text = element_text(size = 8),
        axis.title.x = element_blank(),
        strip.background = element_blank(),
        strip.placement = "outside",
        strip.text = element_text(size = 9)) +
  ylab("Diversity(H)")
#> Warning: Removed 25 rows containing missing values (`geom_line()`).
```



```
ggsave(file='figures/shannon_2.png',
       width = 5, height = 2.2)
#> Warning: Removed 25 rows containing missing values (`geom_line()`).
ggsave('figures/shannon_2.pdf', device = cairo_pdf,
       width = 5, height = 2.2)
#> Warning: Removed 25 rows containing missing values (`geom_line()`).
```

Single Species Models

Model Choice

Our model alternatives are similar to the choices we had for the Total Density model. The problem is, we can't use any of the continuous data distributions in GAMS with zero values (at least relying on the canonical link functions) because ($\log(0) = -\text{Inf}$; $1/0 = \text{Inf}$, $1 / 0*0 = \text{Inf}$). The easiest solution is to add some finite small quantity to the density data, and predict that. Here we predict $\log(\text{Density} + 1)$ using Gaussian models.

Automating Analysis of Separate Species

I'm going to automate analysis of all selected species by using a "nested" Tibble. This is a convenient alternative to writing a "for" loop to run multiple identical analyses.

I create a "long" data source, based on the reduced data set that omits low salinity samples.

```
spp_data <- drop_low %>%
  select(Yearf, Season, Station, Temp,
         Sal, Turb, Chl, Fish,
         Acartia, Balanus, Eurytemora) %>%
  pivot_longer(-c(Yearf:Fish), names_to = 'Species', values_to = 'Density')
```

Next, I create a function to run the analysis. This function takes a data frame or tibble as an argument. The tibble must have data columns with the correct names.

The initial model fits for some species had a lot of wiggles in them, to an extent that I thought did not make much scientific sense, so I decided to reduce the dimensionality of the GAM smoothers, by adding the parameter $k = 4$. Lower numbers constrain the GAM to fit smoother lines.

```
my_gam <- function(.dat) {
  gam(log1p(Density) ~
    s(Temp, bs="ts", k = 5) +
    s(Sal, bs="ts", k = 5) +
    s(log(Turb), bs="ts", k = 5) +
    s(log(Chl), bs="ts", k = 5) +
    s(log1p(Fish), bs="ts", k = 5) +
    s(Yearf, bs = 're'),
    data = .dat, family = "gaussian")
}
```

Next, I create the nested tibble, and conduct the analysis on each species...

```
spp_analysis <- spp_data %>%
  group_by(Species) %>%
  nest() %>%
  mutate(gam_mods = map(data, my_gam))
```

And finally, output the model results. I can do that in a "for" loop, but it's Awkward to look through a long list of output, so I step through each species in turn.

Acartia

```
spp = 'Acartia'
mod <- spp_analysis$gam_mods[spp_analysis$Species == spp][[1]]
dat <- spp_analysis$data[spp_analysis$Species == spp][[1]]
summary(mod)
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> log1p(Density) ~ s(Temp, bs = "ts", k = 5) + s(Sal, bs = "ts",
#>      k = 5) + s(log(Turb), bs = "ts", k = 5) + s(log(Chl), bs = "ts",
#>      k = 5) + s(log1p(Fish), bs = "ts", k = 5) + s(Yearf, bs = "re")
#>
#> Parametric coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    6.598      0.371   17.78  <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:
#>              edf Ref.df      F  p-value
#> s(Temp)        3.6631     4 31.950 < 2e-16 ***
#> s(Sal)          3.2713     4  7.570 0.000232 ***
#> s(log(Turb))    0.6538     4  0.637 0.076037 .
#> s(log(Chl))     0.7323     4  1.316 0.055331 .
#> s(log1p(Fish))  0.7316     4  0.610 0.080622 .
#> s(Yearf)        3.5153     4 11.237 6.14e-07 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) =  0.763   Deviance explained = 81.8%
#> GCV = 0.93657   Scale est. = 0.70553   n = 55
```

Combined Graphic

Generate Separate Marginal Means

```
temp <- marginal_analysis(dat, Temp, mod, .logx = FALSE)
```

```
sal <- marginal_analysis(dat, Sal, mod, .logx = FALSE)
```

```
names(temp) <- c("Pred", names(temp)[2:6])
names(sal) <- c("Pred", names(sal)[2:6])
```

Name match

Assemble data

```
fancy_temp <- expression("Temperature (" * degree * "C)")
fancy_sal  <- expression("Salinity" ~ "(PSU)")

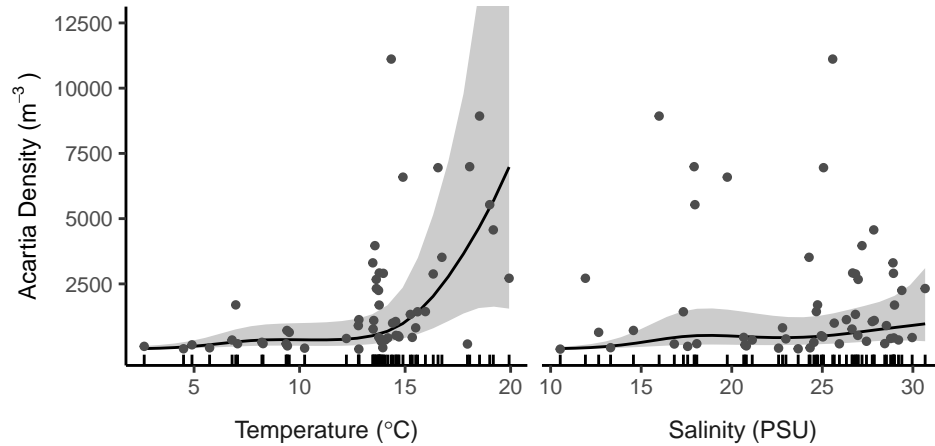
dat <- dat %>%
  select(Temp, Sal, Density) %>%
  pivot_longer(c(Temp, Sal), names_to = "source", values_to = "Pred") %>%
  mutate(source = factor(source,
                        levels = c("Temp", "Sal"),
                        labels = c(fancy_temp, fancy_sal)))

emms <- bind_rows(Temperature = temp, Salinity = sal, .id = "source") %>%
  mutate(source = factor(source,
                        levels = c("Temperature", "Salinity"),
                        labels = c(fancy_temp, fancy_sal)))
```

Generate Plot

```
ggplot(emms, aes(Pred, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = dat, mapping = aes(x = Pred, y = Density),
            size = 1, color = "grey30") +
  geom_rug(data = dat, mapping = aes(x = Pred, y = NULL)) +
  facet_wrap(~source, scales = "free_x", strip.position = "bottom",
            labeller=label_parsed) +

  theme(axis.title = element_text(size = 9),
        axis.text = element_text(size = 8),
        axis.title.x = element_blank(),
        strip.background = element_blank(),
        strip.placement = "outside",
        strip.text = element_text(size = 9)) +
  ylab(expression("Acartia Density (" * m ^{-3} ~ ")") ) +
  scale_y_continuous(breaks = c(1:5*2500)) +
  coord_cartesian(ylim = c(0, 12500))
```

```
ggsave(file='figures/Acartia_2.png',
        width = 5.04, height = 2.5)
ggsave('figures/Acartia_2.pdf', device = cairo_pdf,
        width = 5.04, height = 2.5)
```

Balanus

```
spp = 'Balanus'
mod <- spp_analysis$gam_mods[spp_analysis$Species == spp][[1]]
dat <- spp_analysis$data[spp_analysis$Species == spp][[1]]
summary(mod)
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> log1p(Density) ~ s(Temp, bs = "ts", k = 5) + s(Sal, bs = "ts",
#>          k = 5) + s(log(Turb), bs = "ts", k = 5) + s(log(Chl), bs = "ts",
#>          k = 5) + s(log1p(Fish), bs = "ts", k = 5) + s(Yearf, bs = "re")
#>
#> Parametric coefficients:
#>               Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   3.6930      0.6478   5.701 8.74e-07 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:
#>               edf Ref.df      F  p-value
#> s(Temp)        9.192e-01     4  2.998  0.00414 **
#> s(Sal)         1.782e-10     4  0.000  0.52552
#> s(log(Turb))    1.967e+00     4  1.779  0.06016 .
#> s(log(Chl))     1.004e+00     4 14.125 2.07e-05 ***
#> s(log1p(Fish))  1.686e+00     4  0.691  0.22444
#> s(Yearf)        3.568e+00     4  7.912 1.75e-05 ***
#> ---
```

```
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) =  0.581   Deviance explained = 65.2%
#> GCV = 2.7021   Scale est. = 2.2038       n = 55
```

Generate Separate Marginal Means

```
temp <- marginal_analysis(dat, Temp, mod, .logx = FALSE)
```

```
chl <- marginal_analysis(dat, Chl, mod, .logx = TRUE)
```

```
names(temp) <- c("Pred", names(temp)[2:6])
names(chl) <- c("log(Pred)", names(chl)[2:6], "Pred")
chl <- chl[,c(7, 2:6)]
```

Name match

Assemble data

```
fancy_temp <- expression("Temperature (" * degree * "C)")
fancy_chl <- expression("Chlorophyll (" * mu * g * L ^-1 ~")")

dat <- dat %>%
  select(Temp, Chl, Density) %>%
  pivot_longer(c(Temp, Chl), names_to = "source", values_to = "Pred") %>%
  mutate(source = factor(source,
    levels = c("Temp", "Chl"),
    labels = c(fancy_temp, fancy_chl)))

emms <- bind_rows(Temperature = temp, Chlorophyll = chl, .id = "source") %>%
  mutate(source = factor(source,
    levels = c("Temperature", "Chlorophyll"),
    labels = c(fancy_temp, fancy_chl)))
```

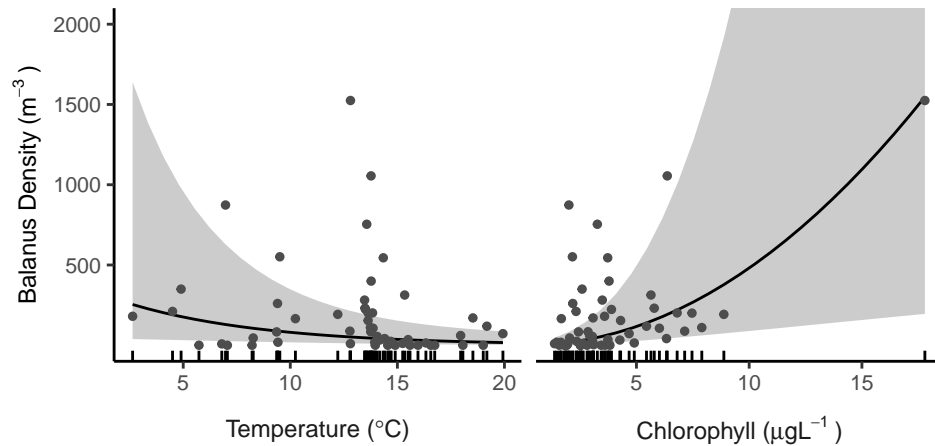
Generate Plot

```
ggplot(emms, aes(Pred, response)) +
  geom_ribbon(aes(ymin = lower.CL, ymax = upper.CL), fill = "grey80") +
  geom_line() +
  geom_point(data = dat, mapping = aes(x = Pred, y = Density),
    size = 1, color = "grey30") +
  geom_rug(data = dat, mapping = aes(x = Pred, y = NULL)) +
  facet_wrap(~source, scales = "free_x", strip.position = "bottom",
    labeller=label_parsed) +
```

```

theme(axis.title = element_text(size = 9),
      axis.text = element_text(size = 8),
      axis.title.x = element_blank(),
      strip.background = element_blank(),
      strip.placement = "outside",
      strip.text = element_text(size = 9)) +
ylab(expression("Balanus Density (" * m ^{-3} ~ ")")) +
scale_y_continuous(breaks = c(1:4*500)) +
coord_cartesian(ylim = c(0, 2000))

```



```

ggsave(file='figures/Balanus_2.png',
        width = 5.04, height = 2.5)
ggsave('figures/Balanus_2.pdf', device = cairo_pdf,
        width = 5.05, height = 2.5)

```